

The C10K Problem and Solutions

References:

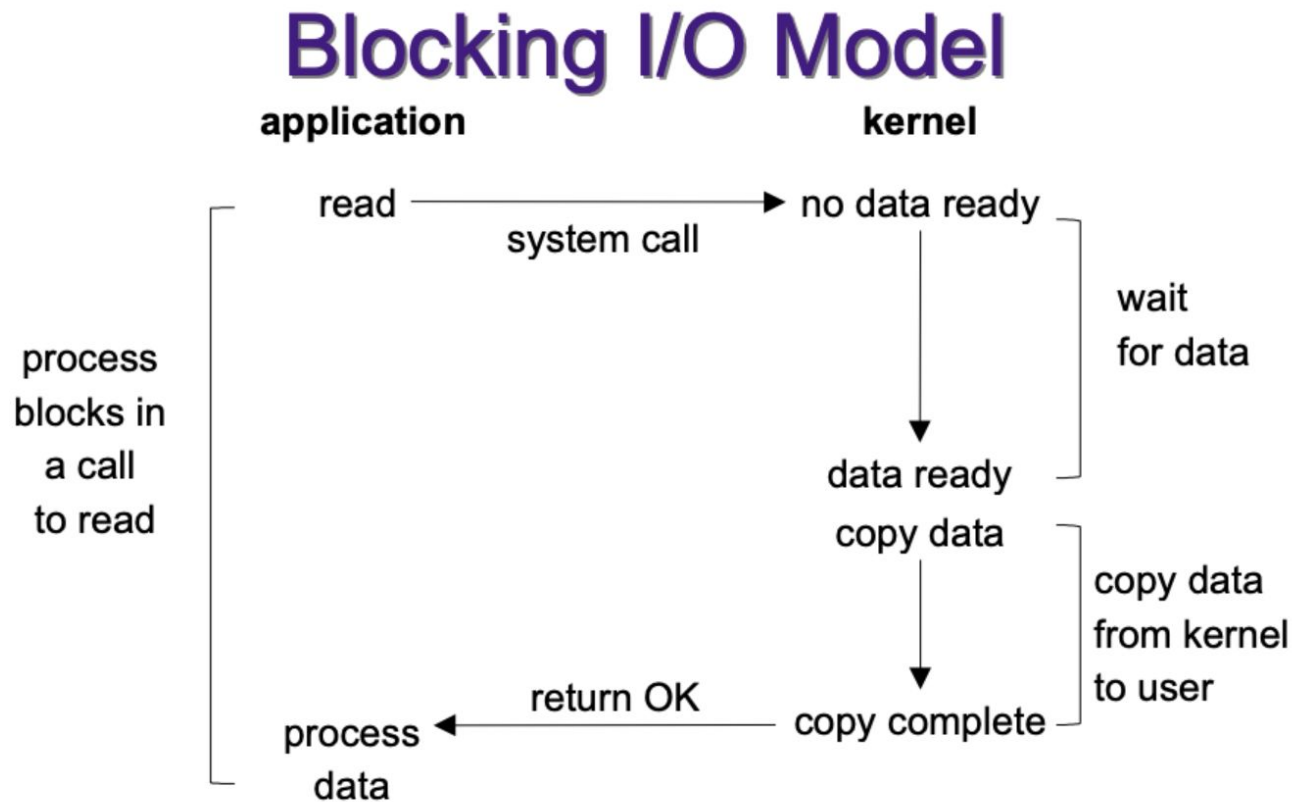
- <http://www.kegel.com/c10k.html>
- <http://www.monkey.org/~provos/libevent>
- <http://byteliu.com/2019/05/08/LINUX-%E2%80%93-IO-MULTIPLEXING-%E2%80%93-SELECT-VS-POLL-VS-EPOLL/>
- A Scalable and Explicit Event Delivery Mechanism for UNIX,
http://static.usenix.org/event/usenix99/full_papers/banga/banga.pdf
- Acknowledgement: Modified from the slides of Che-Yi Lin and Hao-Yun Liu.

Outline

- Background
 - Types of I/O
 - Problem of select()
- The C10K problem
- Asynchronous I/O (AIO)
- Design of networking software
- Using libevent
- Benchmarks

Background: Types of I/O

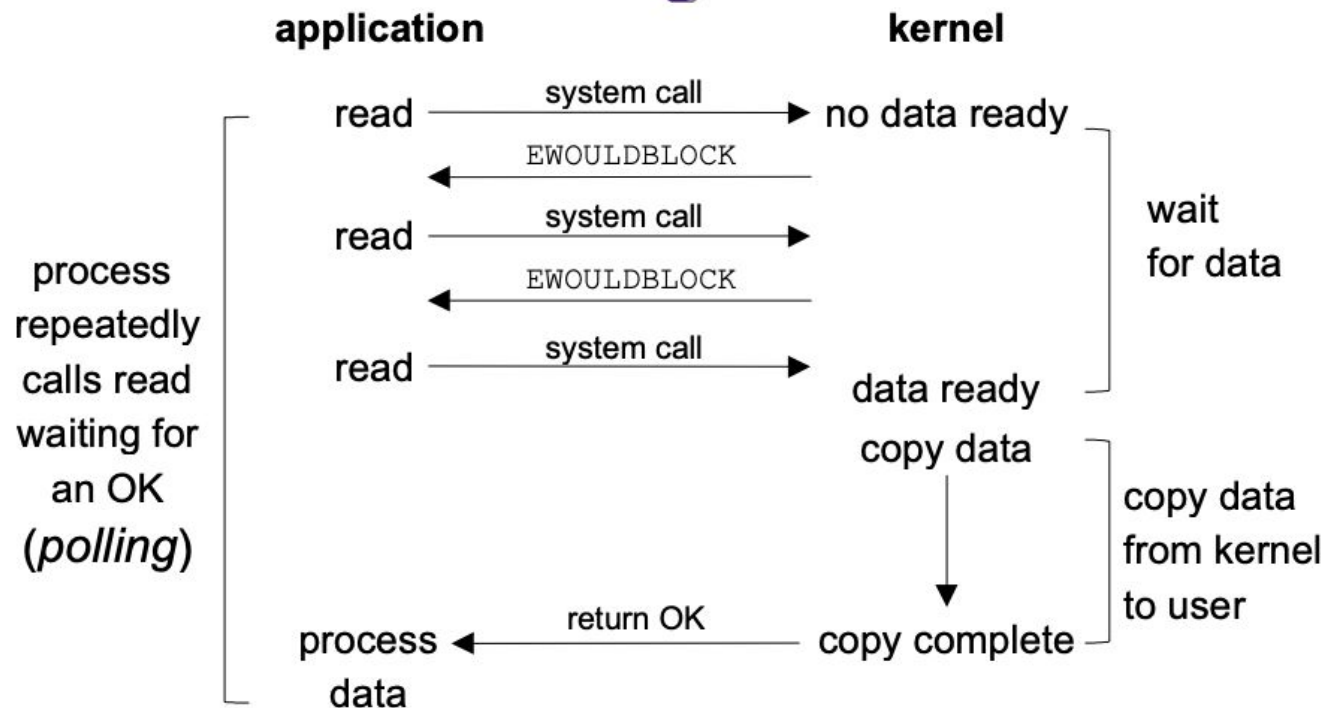
- Wait until data is ready



Background: Types of I/O (cont.)

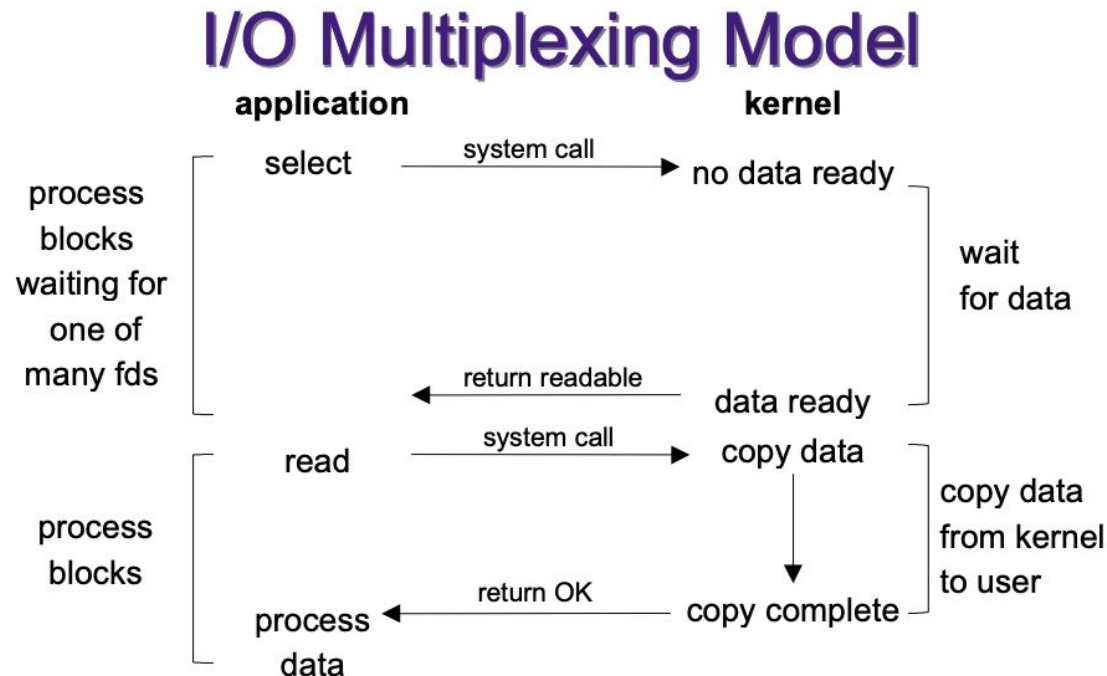
- Return an error if data was not ready

Nonblocking I/O Model



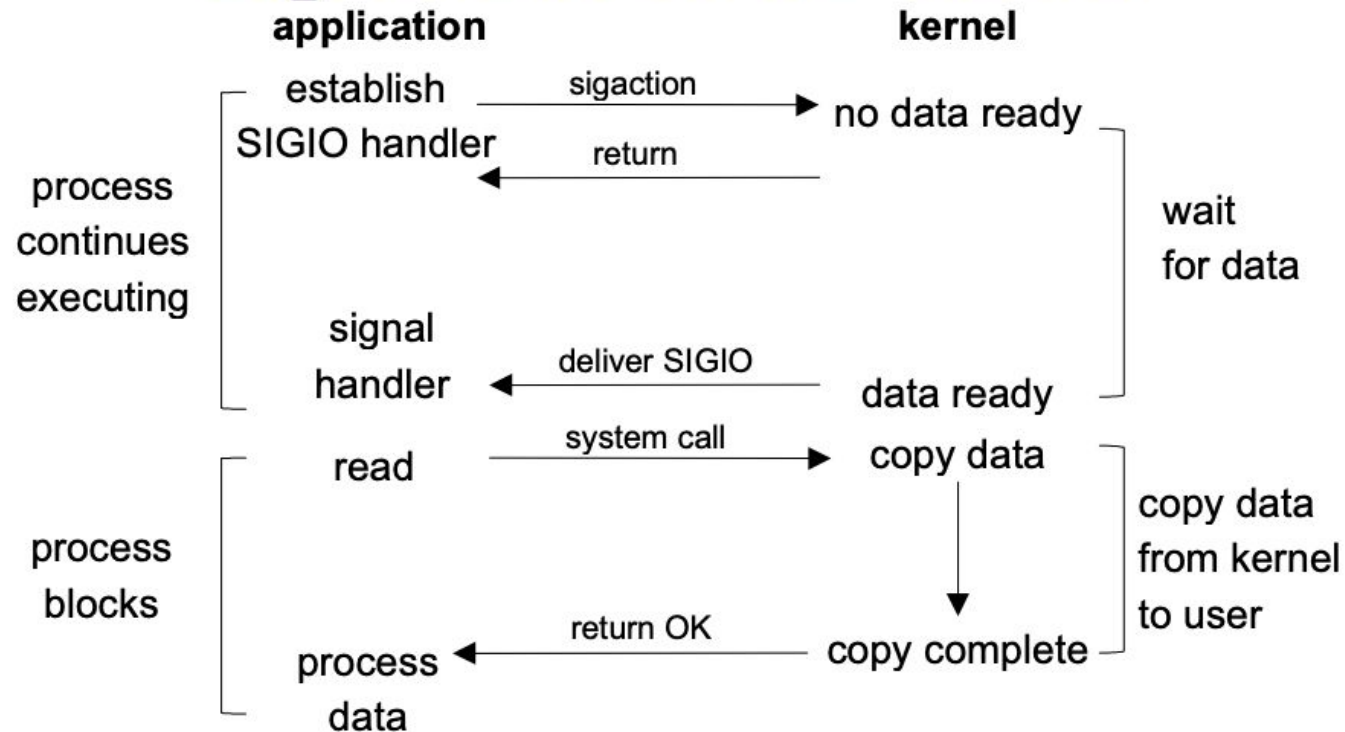
Background: Types of I/O (cont.)

- Block until the given file descriptors are ready to perform I/O
- Block `select()` or `poll()`, instead of blocking I/O
 - `select()`
 - `poll()`
 - `epoll()`



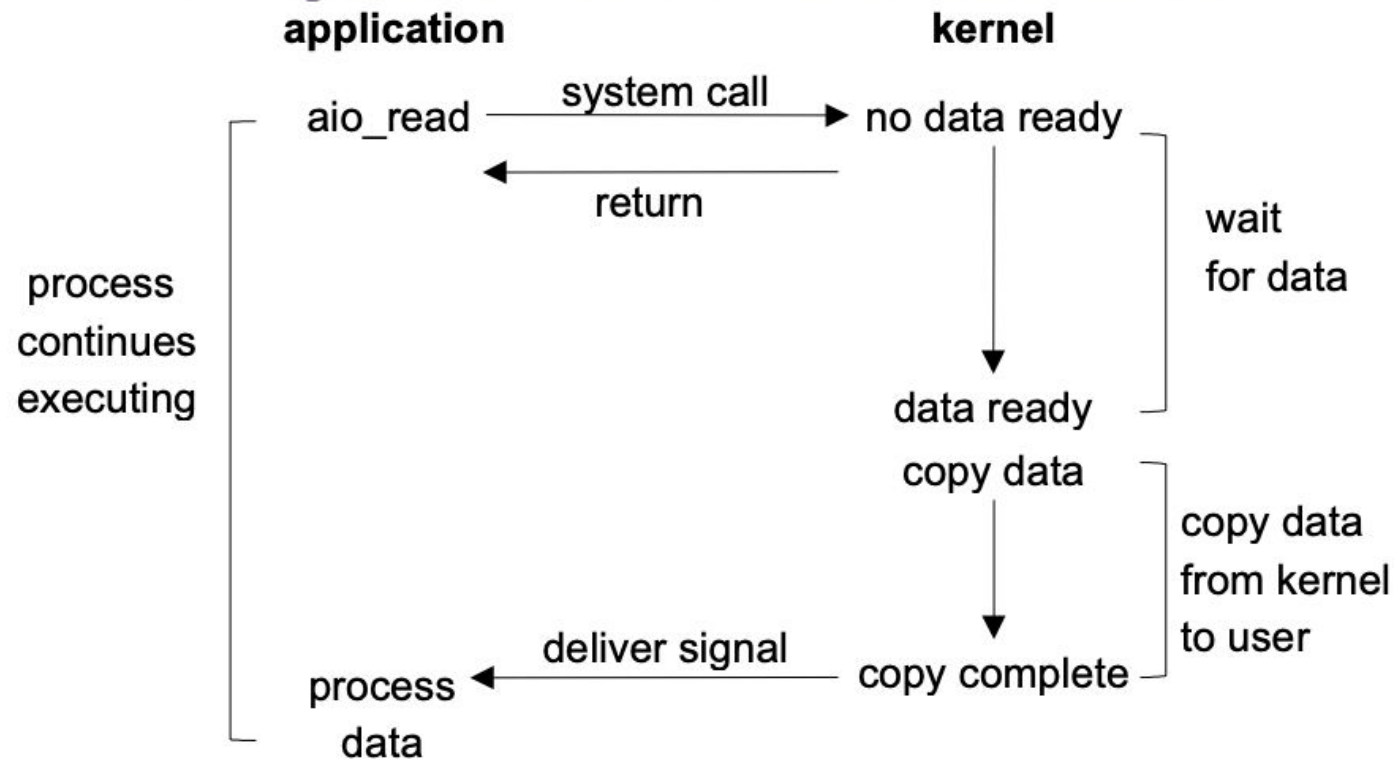
Background: Types of I/O (cont.)

Signal Driven I/O Model

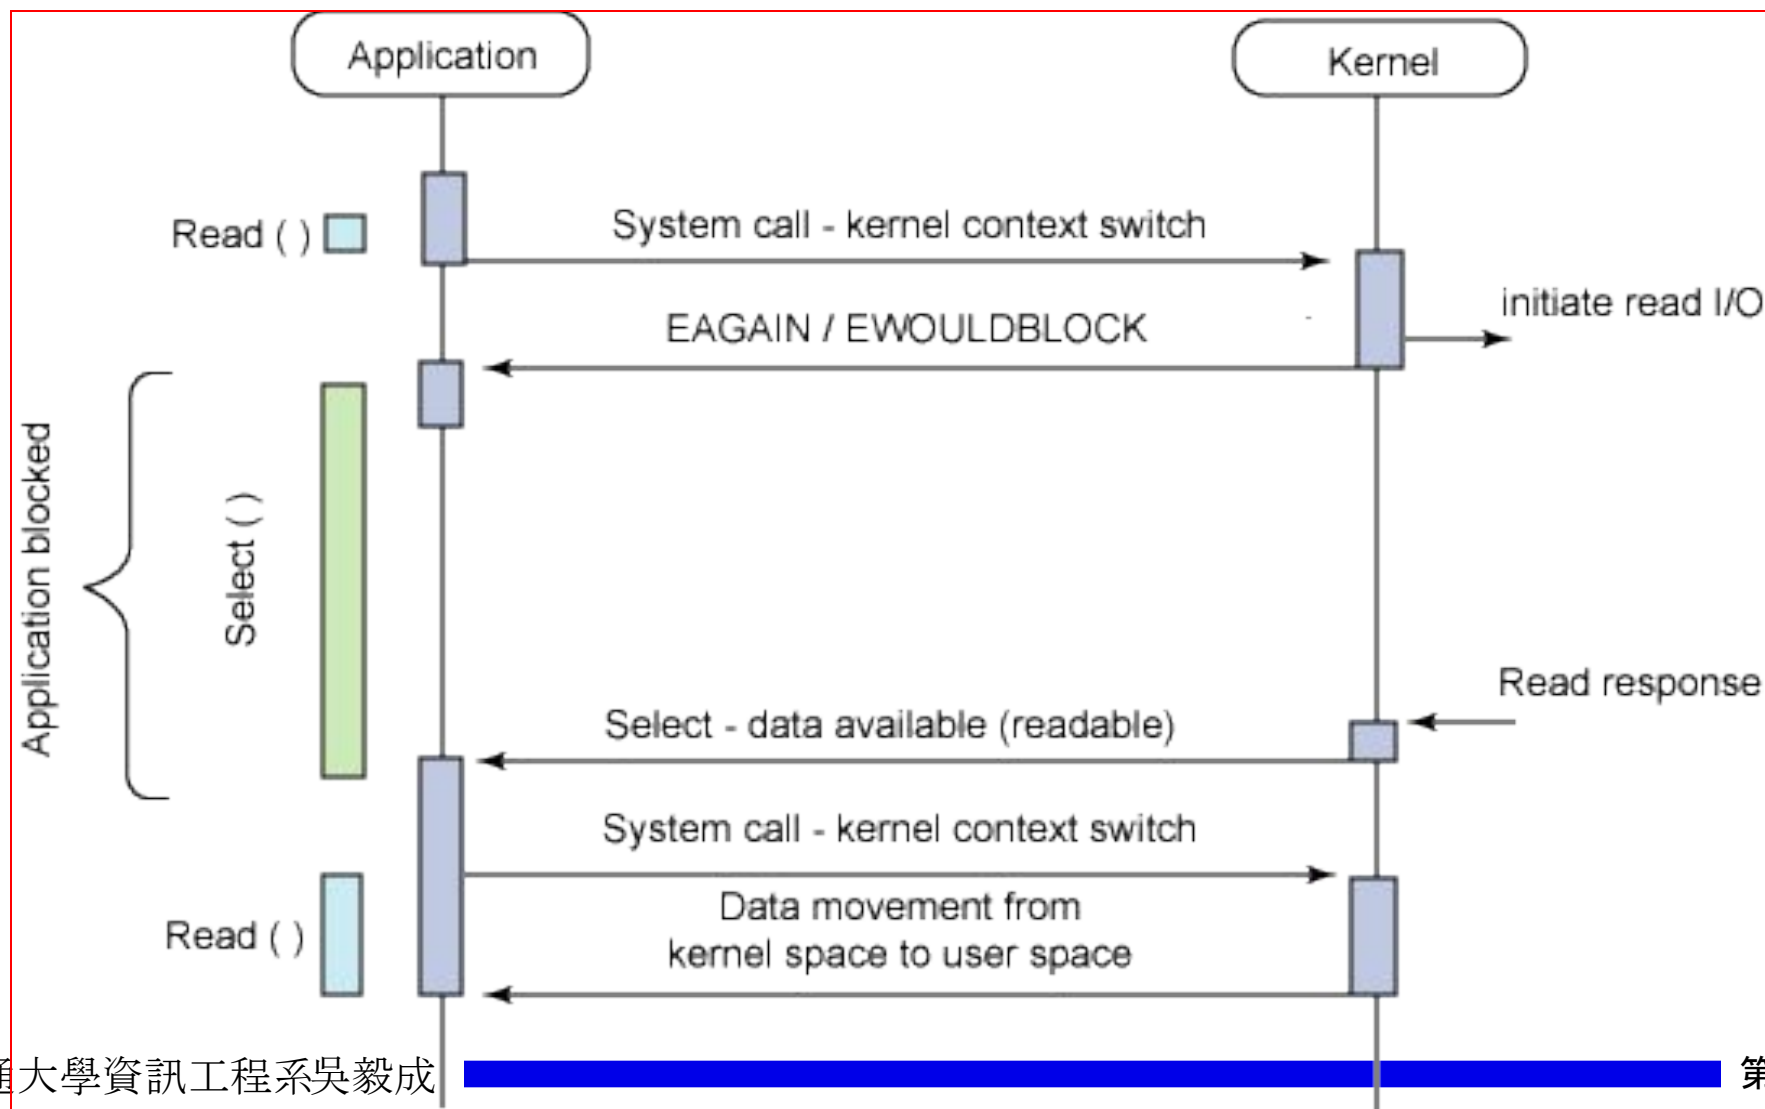


Background: Types of I/O (cont.)

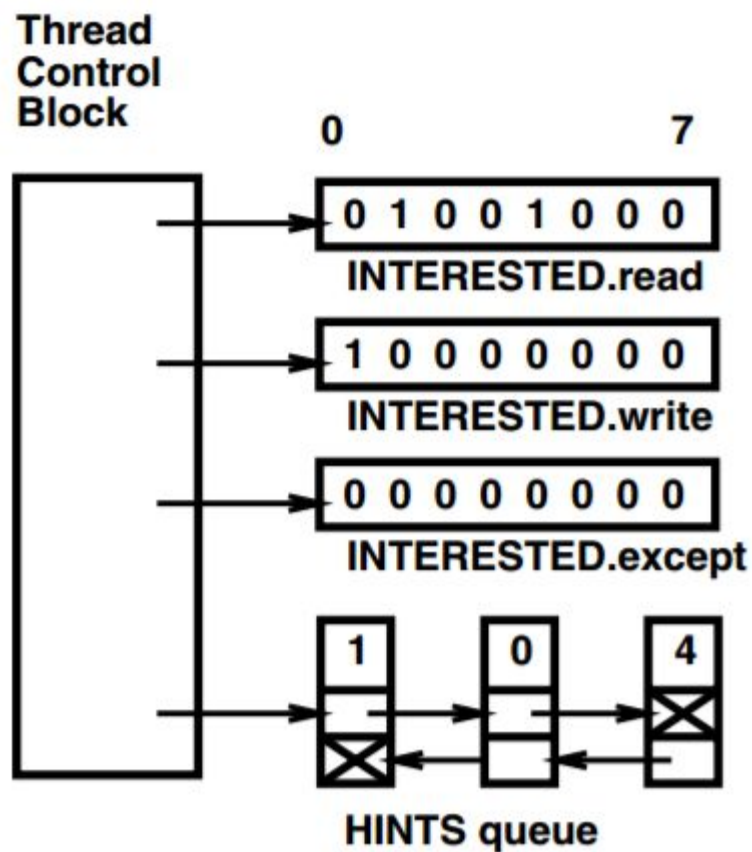
Asynchronous I/O Model



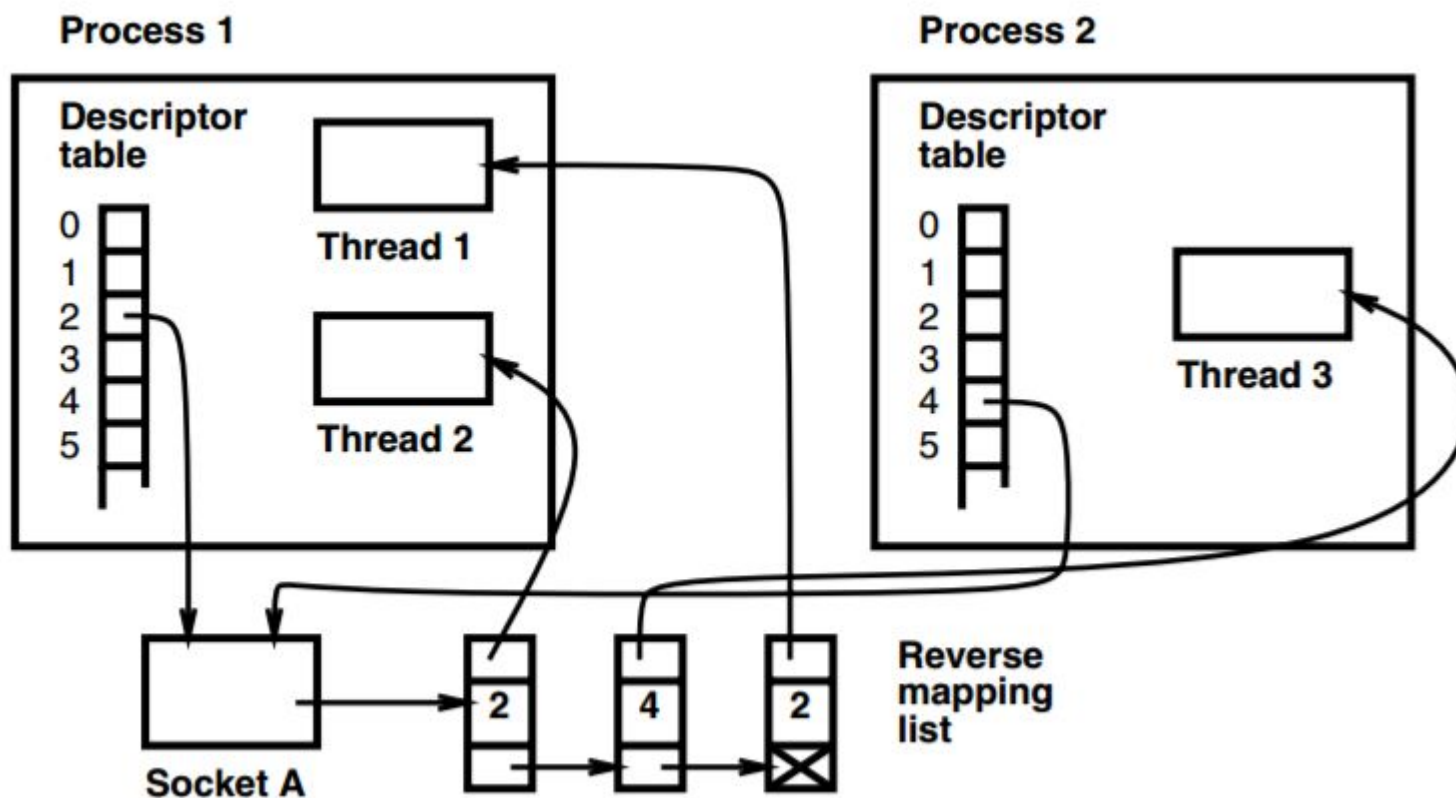
Background: select()



Per-Thread Data Structure



Per-Socket Data Structure

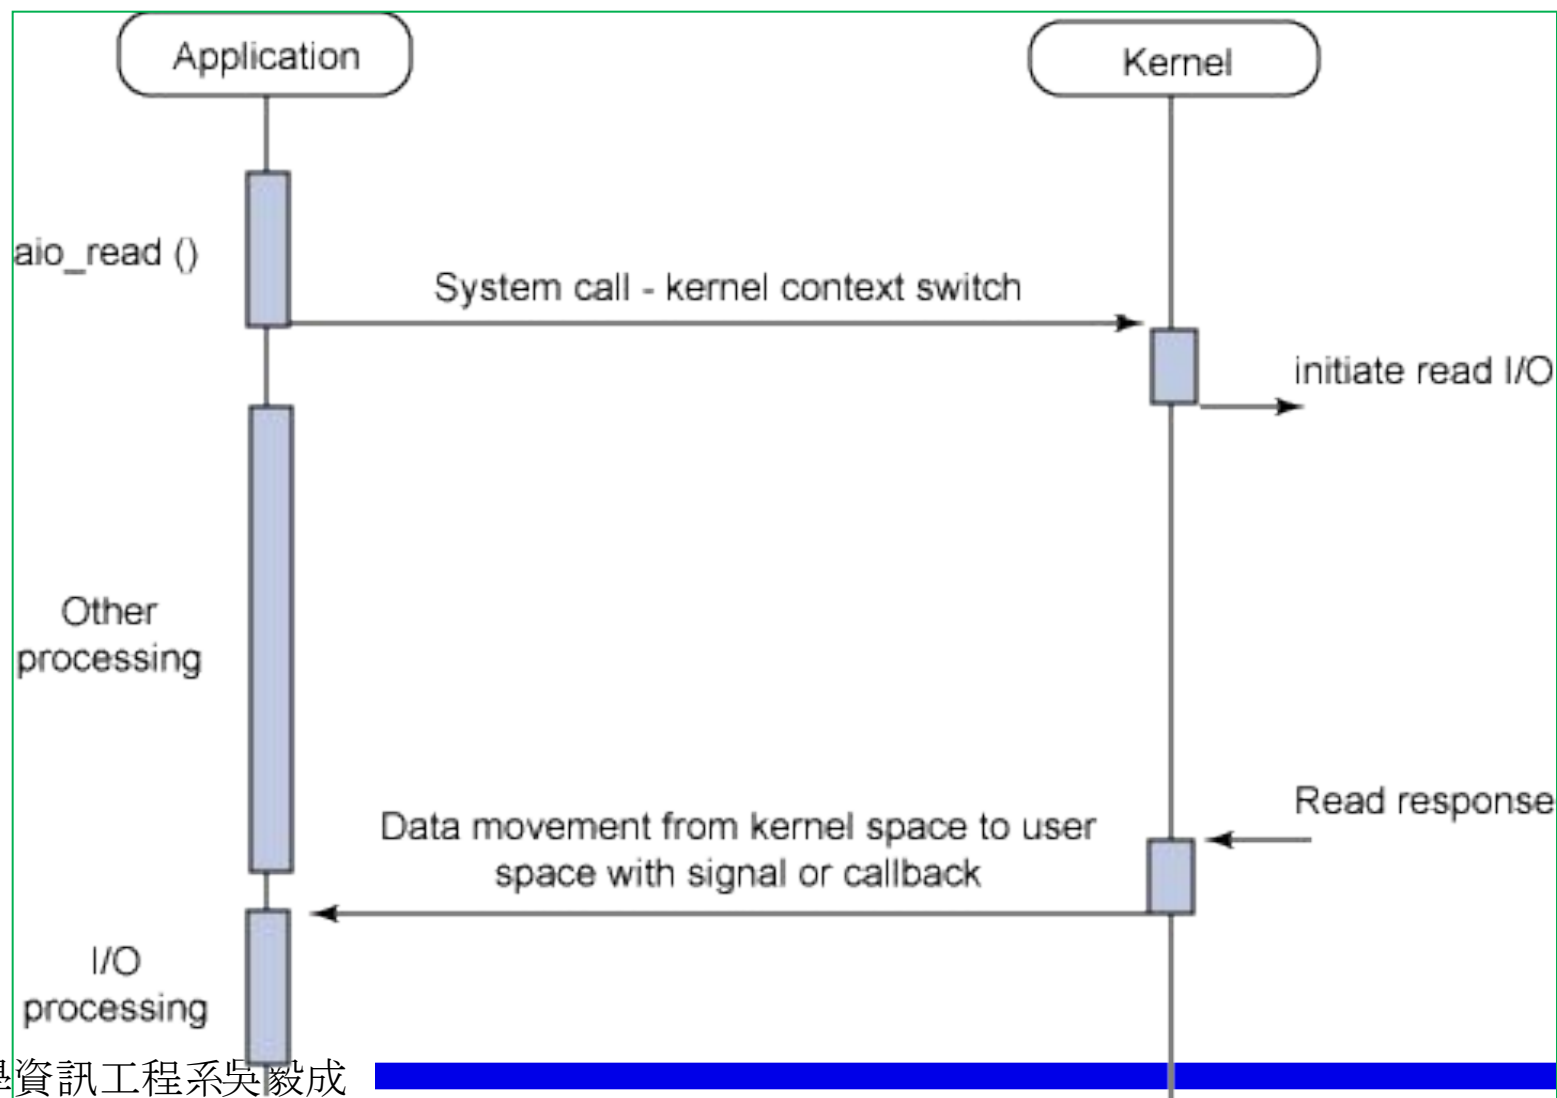


The C10K problem

- Web servers have to handle ten thousand clients simultaneously
- Web is a big place now
- Hardware is no longer the bottleneck

<http://www.kegel.com/c10k.html>

Asynchronous I/O (AIO)



AIO in Linux

Blocking

Non-blocking

Synchronous

Read/write

Read/write
(O_NONBLOCK)

Asynchronous

i/O multiplexing
(select/poll)

AIO

AIO in Linux (cont.)

- Introduced in Linux kernel 2.6 (released at 2008) and also available in 2.4 if patched.
- The completion of I/O can be notified by two method.
 - Signal.
 - Register a completion handler function to create a new thread.
- API:
 - aio_read
 - aio_error
 - aio_return
 - aio_write
 - aio_suspend
 - aio_cancel
 - lio_listio

Design of application

- Use **fork()**

defect: High overhead for each connection

solution: Return to accept() and child process die automatically

example: Apache 1.3

- Use **pthread_create()**

defect: Thread-safe and Memory-leak problems

solution: Use Thread-safe library and Garbage collection library

example: Apache 2.0 Thread MPM

Design of application (cont.)

- Event-based process

advantage: Without overhead of create process or thread, no need to use Share Memory or Mutex for process / thread

hard to implement:

- BSD □ kqueue(), Linux □ epoll(), Solaris □ /dev/poll
None of these are Standard !
- Buffering of nonblocking I/O

Solution: libevent library

Using libevent

- **libevent** by Niels Provos
<http://www.monkey.org/~provos/libevent/>
- Is a lightweight C I/O framework
- Support **kqueue()**, **epoll()**, **/dev/poll**, and the traditional **select()**, **poll()**
- Under 3-clause BSD license !!


Using libevent (cont.)

- Execute a callback function when a specific event occurs on a file descriptor or after a timeout has been reached
- Replace the event loop found in event driven network servers
- An application just needs to call `event_dispatch()` and then add or remove events dynamically without having to change the event loop
- Can also be used for multi-threaded applications

Using libevent (cont.)

```
/* Initial libevent. */  
event_init();  
  
/* Create event. */  
struct event ev;  
event_set(&ev, sfd, EV_READ | EV_PERSIST,  
          connection_accept, &ev);  
  
/* Add event. */  
event_add(&ev, NULL);  
  
event_dispatch();
```

Benchmarks of libevent



Benchmarks (cont.)

